

```
!pip install kaggle
```

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.11/dist-packages (1.7.4.2)
Requirement already satisfied: bleach in /usr/local/lib/python3.11/dist-packages (from kaggle) (6.2.0)
Requirement already satisfied: certifi<=14.05.14 in /usr/local/lib/python3.11/dist-packages (from kaggle) (2025.1.31)
Requirement already satisfied: charset-normalizer in /usr/local/lib/python3.11/dist-packages (from kaggle) (3.4.1)
Requirement already satisfied: idna in /usr/local/lib/python3.11/dist-packages (from kaggle) (3.10)
Requirement already satisfied: protobuf in /usr/local/lib/python3.11/dist-packages (from kaggle) (5.29.4)
Requirement already satisfied: python-dateutil<=2.5.3 in /usr/local/lib/python3.11/dist-packages (from kaggle) (2.8.2)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.11/dist-packages (from kaggle) (8.0.4)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from kaggle) (2.32.3)
Requirement already satisfied: setuptools<=21.0.0 in /usr/local/lib/python3.11/dist-packages (from kaggle) (75.2.0)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.11/dist-packages (from kaggle) (1.17.0)
Requirement already satisfied: text-unidecode in /usr/local/lib/python3.11/dist-packages (from kaggle) (1.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from kaggle) (4.67.1)
Requirement already satisfied: urllib3>=1.15.1 in /usr/local/lib/python3.11/dist-packages (from kaggle) (2.3.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.11/dist-packages (from kaggle) (0.5.1)
```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

```
! mkdir ~/.kaggle
```

```
cp /content/drive/MyDrive/Kaggle_Api/kaggle.json ~/.kaggle/
```

```
! chmod 600 ~/.kaggle/kaggle.json
```

```
! kaggle datasets download arshkon/linkedin-job-postings
```

```
Dataset URL: https://www.kaggle.com/datasets/arshkon/linkedin-job-postings
License(s): CC-BY-SA-4.0
linkedin-job-postings.zip: Skipping, found more recently modified local copy (use --force to force download)
```

```
! unzip linkedin-job-postings.zip
```

```
Archive: linkedin-job-postings.zip
  inflating: companies/companies.csv
  inflating: companies/company_industries.csv
  inflating: companies/company_specialities.csv
  inflating: companies/employee_counts.csv
  inflating: jobs/benefits.csv
  inflating: jobs/job_industries.csv
  inflating: jobs/job_skills.csv
  inflating: jobs/salaries.csv
  inflating: mappings/industries.csv
  inflating: mappings/skills.csv
  inflating: postings.csv
```

## ✓ Step1: Load and Explore Dataset

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from scipy import stats

# Step 1: Load and Explore the Dataset
# For this example, we'll assume the data is loaded from a CSV file
# In a real scenario, replace this with the actual data loading code
print("Step 1: Load and Explore the Dataset")

# Using the sample data provided

# Load the dataset
df = pd.read_csv('postings.csv')

# Get basic information about the dataset
print(f"Dataset shape: {df.shape}")
print("\nData types:")
print(df.dtypes)
```

```
print("\nBasic statistics:")
print(df.describe().T)
print("\nMissing values:")
print(df.isnull().sum())
```

```

closed_time      1073.0   1.712928e+12   3.622893e+08   1.712346e+12
listed_time      123849.0   1.713204e+12   3.989122e+08   1.711317e+12
sponsored        123849.0   0.000000e+00   0.000000e+00   0.000000e+00
normalized_salary  36073.0   2.053270e+05   5.097627e+06   0.000000e+00
zip_code         102977.0   5.040049e+04   3.025223e+04   1.001000e+03
fips             96434.0   2.871388e+04   1.601593e+04   1.003000e+03

```

```

              25%          50%          75%          max
job_id      3.894587e+09   3.901998e+09   3.904707e+09   3.906267e+09
max_salary   4.828000e+01   8.000000e+04   1.400000e+05   1.200000e+08
company_id   1.435200e+04   2.269650e+05   8.047188e+06   1.034730e+08
views        3.000000e+00   4.000000e+00   8.000000e+00   9.975000e+03
med_salary   1.894000e+01   2.550000e+01   2.510500e+03   7.500000e+05
min_salary   3.700000e+01   6.000000e+04   1.000000e+05   8.500000e+07
applies      1.000000e+00   3.000000e+00   8.000000e+00   9.670000e+02
original_listed_time  1.712863e+12   1.713395e+12   1.713478e+12   1.713573e+12
remote_allowed  1.000000e+00   1.000000e+00   1.000000e+00   1.000000e+00
expiry       1.715481e+12   1.716042e+12   1.716088e+12   1.729125e+12
closed_time   1.712670e+12   1.712670e+12   1.713283e+12   1.713562e+12
listed_time   1.712886e+12   1.713408e+12   1.713484e+12   1.713573e+12
sponsored     0.000000e+00   0.000000e+00   0.000000e+00   0.000000e+00
normalized_salary  5.200000e+04   8.150000e+04   1.250000e+05   5.356000e+08
zip_code      2.411200e+04   4.805900e+04   7.820100e+04   9.990100e+04
fips          1.312100e+04   2.918300e+04   4.207700e+04   5.604500e+04

```

```

Missing values:
job_id      0
company_name 1719
title       0
description  7
max_salary  94056
pay_period  87776
location    0
company_id  1717
views       1689
med_salary  117569
min_salary  94056
formatted_work_type  0
applies     100529
original_listed_time  0
remote_allowed  108603
job_posting_url  0
application_url  36665
application_type  0
expiry         0
closed_time    122776
formatted_experience_level  29409
skills_desc    121410
listed_time    0
posting_domain  39968
sponsored      0
work_type      0
currency       87776
compensation_type  87776
normalized_salary  87776
zip_code       20872
fips           27415
dtype: int64

```

## ✓ Step 2: Data Cleaning - Handle Missing Values and Normalize Salary Data

```

print("\n\nStep 2: Data Cleaning")

# Create a new DataFrame to avoid modifying the original
df_clean = df.copy()

# Handle missing company names
df_clean['company_name'].fillna('Unknown Company', inplace=True)

# Normalize salary data - convert hourly wages to yearly
def normalize_salary(row):
    if pd.notna(row['max_salary']):
        if row['pay_period'] == 'HOURLY':
            # Assuming 40 hours per week, 52 weeks per year
            return row['max_salary'] * 40 * 52
        else:
            return row['max_salary']
    return row['normalized_salary']

```

```
# If normalized_salary is available, use it; otherwise calculate from max_salary
df_clean['annual_salary'] = df_clean.apply(normalize_salary, axis=1)

print("Normalized salary data:")
print(df_clean[['title', 'max_salary', 'pay_period', 'normalized_salary', 'annual_salary']].head())
```



Step 2: Data Cleaning

<ipython-input-12-69d59120ef21>:7: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col]

```
df_clean['company_name'].fillna('Unknown Company', inplace=True)
Normalized salary data:
```

	title	max_salary	pay_period	\
0	Marketing Coordinator	20.0	HOURLY	
1	Mental Health Therapist/Counselor	50.0	HOURLY	
2	Assitant Restaurant Manager	65000.0	YEARLY	
3	Senior Elder Law / Trusts and Estates Associat...	175000.0	YEARLY	
4	Service Technician	80000.0	YEARLY	

  

	normalized_salary	annual_salary
0	38480.0	41600.0
1	83200.0	104000.0
2	55000.0	65000.0
3	157500.0	175000.0
4	70000.0	80000.0

### Step 3: Feature Engineering

```
print("\n\nStep 3: Feature Engineering")

# Extract state from location
df_clean['state'] = df_clean['location'].str.extract(r',\s*(\w{2})$')

# Convert listed_time from epoch to datetime
df_clean['listed_date'] = pd.to_datetime(df_clean['listed_time'], unit='ms')

# Create a binary flag for remote jobs
df_clean['is_remote'] = df_clean['remote_allowed'].apply(lambda x: 1 if x == 1.0 else 0)

# Calculate post age in days (from listing to current date)
current_date = pd.Timestamp('2025-04-10') # Using the current date from the prompt
df_clean['post_age_days'] = (current_date - df_clean['listed_date']).dt.days

print("Feature engineering results:")
print(df_clean[['location', 'state', 'listed_time', 'listed_date', 'remote_allowed', 'is_remote', 'post_age_days']].head())
```



Step 3: Feature Engineering  
Feature engineering results:

	location	state	listed_time	listed_date	remote_allowed	\
0	Princeton, NJ	NJ	1.713398e+12	2024-04-17 23:45:08	NaN	
1	Fort Collins, CO	CO	1.712858e+12	2024-04-11 17:51:27	NaN	
2	Cincinnati, OH	OH	1.713278e+12	2024-04-16 14:26:54	NaN	
3	New Hyde Park, NY	NY	1.712896e+12	2024-04-12 04:23:32	NaN	
4	Burlington, IA	IA	1.713452e+12	2024-04-18 14:52:23	NaN	

  

	is_remote	post_age_days
0	0	357
1	0	363
2	0	358
3	0	362
4	0	356

### Step 4: Exploratory Data Analysis - Salary Distribution

```
print("\n\nStep 4: Exploratory Data Analysis - Salary Distribution")

# Basic statistics for annual salary
print("Annual salary statistics:")
print(df_clean['annual_salary'].describe())

# Create a histogram of annual salaries
```

```
plt.figure(figsize=(10, 6))
sns.histplot(df_clean['annual_salary'].dropna(), kde=True)
plt.title('Distribution of Annual Salaries')
plt.xlabel('Annual Salary (USD)')
plt.ylabel('Frequency')
plt.savefig('salary_distribution.png')
plt.close() # Close the figure to avoid displaying in notebook

print("Salary distribution analysis complete. Histogram would show the distribution pattern.")
```



```
Step 4: Exploratory Data Analysis - Salary Distribution
Annual salary statistics:
count    3.607300e+04
mean     2.275688e+05
std      5.540858e+06
min      0.000000e+00
25%      5.453760e+04
50%      9.000000e+04
75%      1.400000e+05
max      5.720000e+08
Name: annual_salary, dtype: float64
Salary distribution analysis complete. Histogram would show the distribution pattern.
```

## ✓ Step 5: Exploratory Data Analysis - Job Engagement Metrics

```
print("\n\nStep 5: Exploratory Data Analysis - Job Engagement Metrics")

# Calculate view-to-application ratio (where data is available)
df_clean['view_apply_ratio'] = df_clean['applies'] / df_clean['views']

# Calculate the correlation between salary and views
salary_views_corr = df_clean['annual_salary'].corr(df_clean['views'])
print(f"Correlation between annual salary and views: {salary_views_corr:.2f}")

# Scatter plot of salary vs. views
plt.figure(figsize=(10, 6))
sns.scatterplot(x='annual_salary', y='views', data=df_clean)
plt.title('Job Views vs. Annual Salary')
plt.xlabel('Annual Salary (USD)')
plt.ylabel('Number of Views')
plt.savefig('salary_vs_views.png')
plt.close()

print("Job engagement analysis complete. Scatter plot would show relationship between salary and views.")
```



```
Step 5: Exploratory Data Analysis - Job Engagement Metrics
Correlation between annual salary and views: -0.00
Job engagement analysis complete. Scatter plot would show relationship between salary and views.
```

## ✓ Step 6: Geographic Analysis

```
print("\n\nStep 6: Geographic Analysis")

# Count jobs by state
state_counts = df_clean['state'].value_counts()
print("Job count by state:")
print(state_counts)

# Calculate average salary by state
avg_salary_by_state = df_clean.groupby('state')['annual_salary'].mean().sort_values(ascending=False)
print("\nAverage annual salary by state:")
print(avg_salary_by_state)

# Create a bar chart of average salaries by state
plt.figure(figsize=(12, 6))
avg_salary_by_state.plot(kind='bar')
plt.title('Average Annual Salary by State')
plt.xlabel('State')
plt.ylabel('Average Annual Salary (USD)')
plt.xticks(rotation=45)
plt.savefig('salary_by_state.png')
plt.close()
```

```
print("Geographic analysis complete. Bar chart would show salary differences by state.")
```



Step 6: Geographic Analysis

Job count by state:

state

CA	11484
TX	10271
NY	6044
FL	5907
NC	4927
IL	4480
PA	4133
VA	3660
MA	3489
OH	3421
GA	3420
NJ	3286
MI	2857
WA	2708
AZ	2507
CO	2318
MD	1974
MO	1922
TN	1885
MN	1849
WI	1849
IN	1808
SC	1539
CT	1191
KY	1179
OR	1177
LA	1106
AL	1004
IA	995
DC	992
UT	968
KS	931
NV	907
OK	794
AR	665
NE	591
NH	559
NM	499
HI	425
WV	416
ID	413
MS	387
ME	377
DE	320
RI	306
MT	236
ND	235
AK	206
VT	181
SD	165
WY	125
ON	1
QC	1

## ✓ Step 7: Salary Prediction Model (Basic)

```
print("\n\nStep 7: Salary Prediction Model (Basic)")
```

```
# For this simple example, we'll create a basic linear regression model
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import mean_squared_error, r2_score
```

```
# Select features for the model
```

```
features = ['views', 'post_age_days', 'is_remote']
```

```
available_features = [f for f in features if f in df_clean.columns]
```

```
# Prepare the data (drop rows with missing values)
```

```
model_data = df_clean.dropna(subset=['annual_salary'] + available_features)
```

```
if len(model_data) >= 5: # Need at least 5 rows for meaningful split
```

```
    # Split the data into training and testing sets
```

```
    X = model_data[available_features]
```

```
    y = model_data['annual_salary']
```

```
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
    # Train the model
```

```

model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Model trained with features: {available_features}")
print(f"Mean Squared Error: {mse:.2f}")
print(f"R² Score: {r2:.2f}")

# Feature importance
feature_importance = pd.DataFrame({
    'Feature': available_features,
    'Importance': model.coef_
})
print("\nFeature importance:")
print(feature_importance.sort_values('Importance', ascending=False))
else:
    print("Not enough complete data rows to build a prediction model.")
    print("In a real analysis, you would need more data points.")

```



```

Step 7: Salary Prediction Model (Basic)
Model trained with features: ['views', 'post_age_days', 'is_remote']
Mean Squared Error: 7656380123681.46
R² Score: -0.00

Feature importance:
   Feature  Importance
2  is_remote  74075.208572
0    views   -110.530881
1 post_age_days -1192.381103

```

## ✓ Step 8: Text Analysis of Job Descriptions

```

print("\n\nStep 8: Text Analysis of Job Descriptions")

# Basic text analysis
def count_words(text):
    if pd.isna(text):
        return 0
    return len(str(text).split())

df_clean['description_word_count'] = df_clean['description'].apply(count_words)
df_clean['skills_word_count'] = df_clean['skills_desc'].apply(count_words)

print("Job description statistics:")
print(df_clean[['description_word_count', 'skills_word_count']].describe())

# Example of text processing for keyword extraction
from collections import Counter
import re

def extract_keywords(text_series):
    # Combine all text
    all_text = ' '.join(text_series.dropna().astype(str))

    # Clean text and tokenize
    words = re.findall(r'\b[a-zA-Z]{3,}\b', all_text.lower())

    # Remove common stopwords (simplified list)
    stopwords = {'and', 'the', 'for', 'with', 'are', 'this', 'that', 'you', 'our', 'has', 'have'}
    words = [word for word in words if word not in stopwords]

    # Count word frequencies
    word_counts = Counter(words)
    return word_counts.most_common(10)

top_description_keywords = extract_keywords(df_clean['description'])
print("\nTop keywords in job descriptions:")
print(top_description_keywords)

top_skills_keywords = extract_keywords(df_clean['skills_desc'])
print("\nTop keywords in skills requirements:")

```

```
print(top_skills_keywords)
```



Step 8: Text Analysis of Job Descriptions

Job description statistics:

	description_word_count	skills_word_count
count	123849.000000	123849.000000
mean	523.027929	0.500166
std	301.940688	11.201057
min	0.000000	0.000000
25%	298.000000	0.000000
50%	477.000000	0.000000
75%	696.000000	0.000000
max	3400.000000	529.000000

Top keywords in job descriptions:

['experience', 374182), ('work', 349453), ('will', 278659), ('all', 274650), ('team', 250455), ('your', 204374), ('other', 197064),

Top keywords in skills requirements:

['skills', 879), ('experience', 629), ('position', 527), ('following', 478), ('requires', 473), ('ability', 436), ('work', 380), (



## ✓ Step 9: Job Market Segmentation

```
print("\n\nStep 9: Job Market Segmentation")
```

```
# Prepare data for clustering
```

```
cluster_features = ['annual_salary', 'views', 'remote_allowed']
```

```
cluster_data = df_clean[cluster_features].dropna()
```

```
if len(cluster_data) >= 5: # Need at least 5 rows for meaningful clustering
```

```
    # Scale the data
```

```
    scaler = StandardScaler()
```

```
    scaled_data = scaler.fit_transform(cluster_data)
```

```
# Determine optimal number of clusters (using elbow method, simplified for this example)
```

```
max_clusters = min(4, len(cluster_data) - 1) # Simplified for small dataset
```

```
wcss = []
```

```
for i in range(1, max_clusters + 1):
```

```
    kmeans = KMeans(n_clusters=i, random_state=42, n_init=10)
```

```
    kmeans.fit(scaled_data)
```

```
    wcss.append(kmeans.inertia_)
```

```
# Apply K-Means clustering
```

```
k = 2 # Simplified choice for this example
```

```
kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
```

```
cluster_data['cluster'] = kmeans.fit_predict(scaled_data)
```

```
# Analyze clusters
```

```
cluster_analysis = cluster_data.groupby('cluster').mean()
```

```
print("Job market segments by salary and engagement:")
```

```
print(cluster_analysis)
```

```
# Visualize clusters (2D projection)
```

```
plt.figure(figsize=(10, 6))
```

```
sns.scatterplot(x='annual_salary', y='views', hue='cluster', data=cluster_data, palette='viridis')
```

```
plt.title('Job Market Segments')
```

```
plt.xlabel('Annual Salary (USD)')
```

```
plt.ylabel('Number of Views')
```

```
plt.savefig('job_clusters.png')
```

```
plt.close()
```

```
print("Segmentation analysis complete. Scatter plot would show different job market segments.")
```

```
else:
```

```
    print("Not enough complete data rows for clustering analysis.")
```

```
    print("In a real analysis, you would need more data points.")
```



Step 9: Job Market Segmentation

Job market segments by salary and engagement:

	annual_salary	views	remote_allowed
cluster			
0	1.400893e+05	58.06402	1.0
1	3.120000e+08	22.00000	1.0

Segmentation analysis complete. Scatter plot would show different job market segments.

## ✓ Step 10: Insights and Recommendations

```
print("\n\nStep 10: Insights and Recommendations")

# Calculate job market attractiveness score
# Higher score = more competitive compensation and more engagement
df_clean['market_score'] = (
    df_clean['annual_salary'].fillna(df_clean['annual_salary'].median()) / df_clean['annual_salary'].median() * 0.7 +
    df_clean['views'].fillna(df_clean['views'].median()) / df_clean['views'].median() * 0.3
)

# Sort by market score
top_jobs = df_clean.sort_values('market_score', ascending=False).head(3)
print("Top jobs by market attractiveness:")
print(top_jobs[['title', 'company_name', 'annual_salary', 'views', 'market_score']])

# Identify potential undervalued jobs
# Jobs with high view counts but lower salaries
df_clean['salary_percentile'] = df_clean['annual_salary'].rank(pct=True)
df_clean['views_percentile'] = df_clean['views'].rank(pct=True)
df_clean['value_gap'] = df_clean['views_percentile'] - df_clean['salary_percentile']

undervalued_jobs = df_clean.sort_values('value_gap', ascending=False).head(3)
print("\nPotentially undervalued jobs (high interest, lower salary):")
print(undervalued_jobs[['title', 'company_name', 'annual_salary', 'views', 'value_gap']])

# Summary insights
print("\nSummary Insights:")
print(f"1. Average annual salary across all jobs: ${df_clean['annual_salary'].mean():.2f}")
print(f"2. Most jobs are {df_clean['formatted_work_type'].mode()[0]}")
print(f"3. Remote-friendly jobs: {df_clean['is_remote'].sum()} out of {len(df_clean)}")
print(f"4. Average job view count: {df_clean['views'].mean():.1f}")
print(f"5. Most job postings are from {df_clean['state'].mode()[0] if not df_clean['state'].mode().empty else 'various states'}")

print("\nRecommendations:")
print("1. Focus job search on high market score positions for best compensation")
print("2. Consider undervalued jobs for positions with high interest but potentially less competition")
print("3. Research companies in states with higher average salaries")
print("4. For employers: Include comprehensive skills descriptions to attract more qualified candidates")
print("5. For job seekers: Target positions with detailed job descriptions as they tend to offer better compensation")
```



Step 10: Insights and Recommendations

Top jobs by market attractiveness:

	title \
9237	Intellectual Property Associate (246215)
98888	Case Manager RN, Pedi Rheumatology
89082	Cloud Domain Architect

	company_name	annual_salary	views	market_score
9237	Eastridge Workforce Solutions	57200000.0	4.0	4449.188889
98888	Kaiser Permanente	408865600.0	4.0	3180.365778
89082	Applicantz	31200000.0	42.0	2429.816667

Potentially undervalued jobs (high interest, lower salary):

	title	company_name	annual_salary \
46750	UX/UI Designer	MIDIScale	0.0
29777	Human Resources Generalist	Lionsgate	32.0
40894	Remote Software Engineer (Python)	Insight Global	63.0

	views	value_gap
46750	260.0	0.994201
29777	476.0	0.993421
40894	306.0	0.988282

Summary Insights:

1. Average annual salary across all jobs: \$227568.81
2. Most jobs are Full-time
3. Remote-friendly jobs: 15246 out of 123849
4. Average job view count: 14.6
5. Most job postings are from CA

Recommendations:

1. Focus job search on high market score positions for best compensation
2. Consider undervalued jobs for positions with high interest but potentially less competition
3. Research companies in states with higher average salaries
4. For employers: Include comprehensive skills descriptions to attract more qualified candidates
5. For job seekers: Target positions with detailed job descriptions as they tend to offer better compensation



